

CS/ECE 3810: Computer Systems Architecture

Lecture 1: Introduction

Anton Burtsev
September, 2022

Class details

- Undergraduate
 - 203 students
- Instructor: Anton Burtsev

Who am I?

- I build operating systems
 - Since I was your age, i.e., since 2000
- Bits of L4 microkernel, micro-ITRON, XenTT, LCDs, KSplit
 - <https://www.cs.utah.edu/~aburtsev/>

Prospective students

I am looking for students interested in operating systems at all levels from undergraduate to PhD, if you have relevant skills send me an email.

We are building three new operating systems

RedLeaf: a clean-slate operating system in Rust designed to support formal verification of functional correctness ([project web page](#)).

Redshift: a new operating system aimed to support heterogeneous hardware, e.g., FPGAs, GPUs, TPUs, near storage, and near network cores, etc., as first class citizens ([project web page](#)).

Horizon: a new secure hypervisor and secure cloud in which users own their data. Horizon is developed in Rust, and relies on novel techniques of hardware and software isolation, and will implement cloud-wide information flow control ([project web page](#)).

Class details

- Undergraduate
 - 203 students
- Instructor: Anton Burtsev
- Meeting time: 11:50am-1:10pm (Mon/Wed)
- 4 TAs
 - Send us a private message on Piazza
- Web page
 - <https://www.cs.utah.edu/~aburtsev/3810>

More details

- 8-10 small homework assignments
- 1-2 lab-like assignments
 - Requires basics familiarity with UNIX
 - Shell, C
- Midterm
- Final
- Grades are curved
 - Homework: 30%, midterm exam: 30%, final exam: 40% of your grade.
 - You can drop two assignments
 - Late submissions are 0

This course


- Book: Hennessy and Patterson's
 - Computer Organization and Design
- Topics
 - Understanding performance/cost/power
 - Assembly language
 - Computer arithmetic
 - Pipelining
 - Using predictions
 - Memory hierarchies
 - Accelerators
 - Reliability and Security

COMPUTER ORGANIZATION AND DESIGN

THE HARDWARE/SOFTWARE INTERFACE

FIFTH EDITION

DAVID A. PATTERSON
JOHN L. HENNESSY



Course organization

- Lectures
 - High level concepts and abstractions
 - Recorded (and hopefully live stream)
- Reading
 - Hennessy and Patterson
 - Bits of additional notes
- Homeworks
- Exams
 - Open notes (but might include open questions)

Questions?

Why this class?



File Edit Selection View Go Debug Terminal Help main.rs - hello-rust - Visual Studio Code

RUN AND DEBUG Run   

VARIABLES

Locals

- args: { size=2 }
 - [size]: 2
 - [capacity]: 2
 - > [0]: "D:\\CRIME\\hello-rust\\ta..."
 - > [1]: "src/clients.json"
 - > [Raw View]: {buf={ptr={pointer="..."}, e: Variable is optimized away an...}}
- json_data: Variable is optimized away an...

src > main.rs

```
use std::process;
11
12 mod fileworker;
13 mod parser;
14
15 pub use fileworker::read_file;
16 pub use parser::parser;
17
18 type TokioError = std::io::Error;
19
20 #[tokio::main]
21 async fn main() -> Result<(), TokioError> {
22     let args: Vec<String> = env::args().collect();
23     println!("arguments: {:?}", args); You, 8 days ago • file as command
24     if args.len() != 2 {panic!("no file specified");}
25     let file_name = &args[1];
26     let json_data = match fileworker::read_file(&file_name).await {
27         Ok(data) => data,
28         Err(e) => error_message(&file_name, e),
29     };
30     println!("len = {}", json_data.len());
31     println!("{}", json_data);
32     parser::parser(&json_data);
33     Ok(())
34 }
35
```

But can you do it fast?
... or secure

Example 1: Database Join

Main-Memory Hash Joins on Multi-Core CPUs: Tuning to the Underlying Hardware

Cagri Balkesen ^{#1}, Jens Teubner ^{#2}, Gustavo Alonso ^{#3}, M. Tamer Özsü ^{*4}

[#] *Systems Group, Department of Computer Science, ETH Zurich, Switzerland*


^{1, 2, 3} {name.surname}@inf.ethz.ch

^{*} *University of Waterloo, Canada*

⁴ tamer.ozsu@uwaterloo.ca

- <https://15721.courses.cs.cmu.edu/spring2016/papers/balkesen-icde2013.pdf>

Example 1: Database Join



Example 1: Database Join




Fig. 6. Original hash table implementation.




Fig. 7. Our hash table implementation.

Example 2: Virtualization

Virtualization Without Direct Execution or Jitting: Designing a Portable Virtual Machine Infrastructure

Darek Mihowka

Emulators

darekm@emulators.com

Stanislav Shwartsman

Intel Corp.

stanislav.shwartsman@intel.com

- https://bochs.sourceforge.io/Virtualization_Without_Hardware_Final.pdf

Example 2: Virtualization

	1000 MHz Pentium III	2533 MHz Pentium 4	2666 MHz Core 2 Duo
Bochs 2.3.5	882	595	180
Bochs 2.3.6	609	533	157
Bochs 2.3.7	457	236	81

Table 3.1: Windows XP boot time on different hosts

Example 2: Virtualization

3.3 Host branch misprediction as biggest cause of slow emulation performance

Every pipelined processor features branch prediction logic used to predict whether a conditional branch in the instruction flow of a program is likely to be taken or not. Branch predictors are crucial in today's modern, superscalar processors for achieving high performance.

Modern CPU architectures implement a set of sophisticated branch predictions algorithms in order to achieve highest prediction rate, combining both static and dynamic prediction methods. When a branch instruction is executed, the branch history is stored inside the processor. Once branch history is available, the processor can predict branch outcome – whether the branch should be taken and the branch target.

A typical Bochs instruction handler method:

```
void BX_CPU_C::SUB_EdGd(bxInstruction_c *i)
{
    Bit32u op2_32, op1_32, diff_32;


    op2_32 = BX_READ_32BIT_REG(i->nnn());

    if (i->modC0()) {      // reg/reg format
        op1_32 = BX_READ_32BIT_REG(i->rm());
        diff_32 = op1_32 - op2_32;
        BX_WRITE_32BIT_REGZ(i->rm(), diff_32);
    }
    else {                  // mem/reg format
        read_RMW_virtual_dword(i->seg(),
            RMAddr(i), &op1_32);
        diff_32 = op1_32 - op2_32;
        Write_RMW_virtual_dword(diff_32);
    }
    SET.LAZY.FLAGS.SUB32(op1_32, op2_32,
        diff_32);
}
```

Listing 3.1: A typical Bochs instruction handler

Why is it so hard?

Microprocessor Performance




50% improvement every year!!

What contributes to this improvement?

Source: H&P Textbook

What's inside a typical machine?

B360 AORUS Motherboard




CPU



- 1 CPU socket
 - 4 cores
 - 2 logical threads each (hyperthreads)



Hyper-Threading
(logical threads)



Memory



Memory abstraction


$\text{WRITE}(addr, value) \rightarrow \emptyset$

Store *value* in the storage cell identified by *addr*.


$\text{READ}(addr) \rightarrow value$

Return the *value* argument to the most recent WRITE call referencing *addr*.

I/O Devices




Dell R830 4-socket server




Dell Poweredge R830 System Server with 2 sockets
on the main floor and 2 sockets on the expansion

http://www.dell.com/support/manuals/us/en/19/poweredge-r830/r830_om/supported-configurations-for-the-poweredge-r830-system?guid=guid-01303b2b-f884-4435-b4e2-57bec2ce225a&lang=en-us

Multi-socket machines



Dell R830 4-socket server




Dell Poweredge R830 System Server with 2 sockets on the main floor and 2 sockets on the expansion



http://www.dell.com/support/manuals/us/en/19/poweredge-r830/r830_om/supported-configurations-for-the-poweredge-r830-system?guid=guid-01303b2b-f884-4435-b4e2-57bec2ce225a&lang=en-us

Dell R830 Motherboard




What does CPU do internally?



CPU execution loop

- CPU repeatedly reads instructions from memory
- Executes them
- Example


ADD EDX, EAX

// EDX = EAX + EDX





A simple 5-stage pipeline



Thank you!